

Lifetimes

```
error[E0597]: `x` does not live long enough
```

1

The Borrowing Rules

1. No borrower cannot live longer than the owner
 - i.e., no dangling references
2. You may have one or the other of these two kinds of borrows, but not both at the same time:
 - one or more **shared references** (`&T`) to a resource,
 - exactly one **mutable reference** (`&mut T`).
Why have these restrictive rules?
These rules prevent data races.

Data Races

There is a 'data race' when two or more pointers access the same memory location at the same time, where at least one of them is writing, and the operations are not synchronized.

2

Iterator invalidation

- “Iterator invalidation” happens when you try to mutate a collection that you’re iterating over. Rust’s borrow checker prevents this from happening:

```
let mut v = vec![1, 2, 3];
    immutable borrow
for i in &v { occurs here
    mutable borrow attempted here println!("{}{}", i);
    | v.push(34); |
}
    immutable borrow ends here
```

error: cannot borrow `v` as mutable because it is also borrowed as immutable
v.push(34);
^
note: previous borrow of `v` occurs here; the immutable borrow prevents subsequent moves or mutable borrows of `v` until the borrow ends
for i in &v {
 ^
note: previous borrow ends here
for i in &v {
 println!("{}{}", i);
 v.push(34);
}^

- We can’t modify `v` because it’s borrowed by the loop.

3

Use after Free

- References must not live longer than the resource they refer to.** Rust will check the lifetime of your references to ensure that this is true.

```
let y: &i32;
{
    let x = 5;
    y = &x;
}
println!("{}{}", y);
```

error: `x` does not live long enough
y = &x;
^

- `y` is only valid for the scope where `x` exists. As soon as `x` goes away, it becomes invalid to refer to it.

4

Lifetime and Scope

- A variable's lifetime is **a scope in which it is valid**.
- A variable's lifetime **begins when it is created and ends when it is dropped**.
- *A lifetime is a construct the compiler (or more specifically, its borrow checker) uses to ensure all borrows are valid.*
- When we borrow a variable via `&`, the borrow has a lifetime that is
 - determined by where it is declared, and
 - is valid as long as it ends before the lender is destroyed.

*Every reference in Rust has a **lifetime**.*

5

```
// Lifetimes are annotated below with lines denoting the creation
// and destruction of each variable.
// `i` has the longest lifetime because its scope entirely encloses
// both `borrow1` and `borrow2`. The duration of `borrow1` compared
// to `borrow2` is irrelevant since they are disjoint.
fn main() {
    let i = 3; // Lifetime for `i` starts. ——————
    {
        let borrow1 = &i; // `borrow1` lifetime starts. ——————
        println!("borrow1: {borrow1}"); // ——————
    } // `borrow1` ends. ——————
    {
        let borrow2 = &i; // `borrow2` lifetime starts. ——————
        println!("borrow2: {borrow2}"); // ——————
    } // `borrow2` ends. ——————
} // i's lifetime ends. ——————
```

6

Preventing Dangling References with Lifetimes

- The main aim of lifetimes is to prevent *dangling references*.

```
fn main() {  
    let r;  
  
    {  
        let x = 5;  
        r = &x;  
    }  
  
    println!("r: {}", r);  
}  
  
  
  
$ cargo run  
Compiling chapter10 v0.1.0 (file:///projects/chapter10)  
error[E0597]: `x` does not live long enough  
--> src/main.rs:6:13  
  
6         r = &x;  
          ^^^ borrowed value does not live long enough  
7     }  
     - `x` dropped here while still borrowed  
8  
9     println!("r: {}", r);  
          - borrow later used here
```

7

Borrow Checker

- The Rust compiler has a *borrow checker* that compares scopes to determine whether all borrows are valid.

fn main() {

```

let r;
{
    let x = 5;
    r = &x;
}
println!("r: {}", r);
}

```

Big Problem

borrower **owner**

r's lifetime = lifetime(a') > lifetime(b') = x's lifetime

fn main() {

```

let x = 5;
let r = &x;
}
println!("r: {}", r);
}

```

r → ???

r's lifetime(a') < x's lifetime(b')

8

Lifetime Annotation Syntax

- Lifetime annotations describe **the relationships of the lifetimes** of multiple references to each other without affecting the lifetimes.

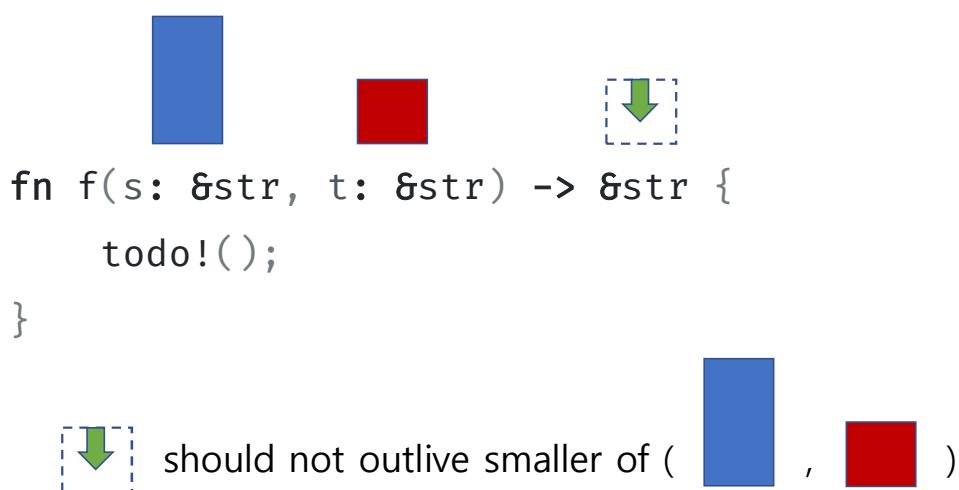
```
&i32          // a reference  
&'a i32       // a reference with an explicit lifetime  
&mut i32     // a mutable reference  
&'a mut i32  // a mutable reference with an explicit lifetime
```

*↑
the lifetime a*

```
fn foo<'a>(...)  
fn foo<'a, 'b>(...)  
fn foo<'a>(x: &'a i32, y: &'a mut i32) -> &'a i32
```

9

What could be wrong?
What should be guaranteed?



10

What is wrong? What should be guaranteed?

```
fn f(s: &str, t: &str) -> &str {  
    todo!();  
}
```

```
error[E0106]: missing lifetime specifier  
--> src/lib.rs:1:27  
|  
1 | fn f(s: &str, t: &str) -> &str {  
|     ----  ----  ^ expected named lifetime parameter  
|  
= help: this function's return type contains a borrowed value, but the signature  
does not say whether it is borrowed from `s` or `t`  
help: consider introducing a named lifetime parameter  
|  
1 | fn f<'a>(s: &'a str, t: &'a str) -> &'a str {  
|     +++++  ++
```

11

What does this lifetime annotation imply?

Hey caller, make sure to **pass only arguments that lives at least as long as the returned reference.**
Otherwise you are doomed.



```
fn f<'a>(s: &'a str, t: &'a str) -> &'a str {  
    todo!();  
}
```

In other words, the returned reference is valid only as long as both parameters are alive.

12

Lifetime Annotations in Function Signatures

- If we want to express the following constraint: “*the returned reference will be valid as long as both the parameters are valid.*”

```
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {  
    if x.len() > y.len() {  
        x  
    } else {  
        y  
    }  
}
```

The lifetime of the reference returned by the “longest()” function is **the same as the smaller of the lifetimes** of the values referred to by the function arguments.

13

Lifetime Annotations in Function Signatures

```
fn demo1() {  
    let s1 = String::from("abcd");  
    let s2 = String::from("xyz");  
  
    let result = longest(&s1, &s2);  
    println!("The longest string is {result}");  
}  
-----  
fn demo2() {  
    let s1 = String::from("abcd");  
  
    let result;  
    {  
        let s2 = String::from("xyz");  
        result = longest(&s1, &s2);  
    }  
    println!("The longest string is {result}");  
}
```

Which one is OK?

14

Thinking in Terms of Lifetimes

- How to specify lifetime parameters depends on what your function is doing.

```
fn first<'a>(x: &'a str, _: &str) -> &'a str {  
    x  
}
```

- When returning a reference from a function, the lifetime parameter for the return type needs to match the lifetime parameter for one of the parameters. Otherwise,

```
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str {  
    let result = String::from("really long string");   
    result.as_str()  
}
```

15

What will happen?

```
fn f<'a>(s: &'a str) -> &'a str { // 'a can be omitted  
    s  
}  
  
let r;  
{  
    let s = String::from("hello");  
    r = f(&s); // by the time `r` is assigned, it is guaranteed  
              // that `s` is still valid  
}  
println!("r: {}", r);
```

error: `s` does not live long enough
label: borrow later used here

16

Multiple lifetimes

- If you have multiple references, you can use the same lifetime multiple times:

```
fn x_or_y<'a>(x: &'a str, y: &'a str) -> &'a str {
```

This says that `x`, `y`, and return value are all alive for the same scope.

- If you wanted `x` and `y` to have different lifetimes, you can use multiple lifetime parameters:

```
fn x_or_y<'a, 'b>(x: &'a str, y: &'b str) -> &'a str {
```

In this example, `x` and `y` have different valid scopes, but the return value has the same lifetime as `x`.

17

Structs with references

```
#[derive(Debug)]
struct City<'a> { // City has lifetime 'a
    name: &'a str, // and name also has lifetime 'a.
    date_founded: u32,
}
```

We need to ensure that any reference to a City cannot outlive the reference to a name it contains.

- It means

"please only take an input for name if it lives at least as long as City".

- It does not mean:

"I will make the input for name live as long as City".

18

Lifetime declaration is mandatory for impl

```
struct Adventurer<'a> {
    name: &'a str,
    hit_points: u32,
}

impl Adventurer {
    fn take_damage(&mut self) {
        self.hit_points -= 20;
        println!("{} has {} hit points left!", self.name, self.hit_points);
    }
}
```

error[E0726]: implicit elided lifetime not allowed here
--> src/lib.rs:5:6
| 5 | impl Adventurer {
| | ^^^^^^^^^^ expected lifetime parameter
| help: indicate the anonymous lifetime
| 5 | impl Adventurer<'_> {
| | +++

19

Lifetime declaration is mandatory for impl

```
struct Adventurer<'a> {
    name: &'a str,
    hit_points: u32,
}

impl<'a> Adventurer<'a> {
    fn take_damage(&mut self) {
        self.hit_points -= 20;
        println!("{} has {} hit points left!", self.name, self.hit_points);
    }
}
```

We repeat '`'a` twice, like on functions: `impl<'a>` defines a lifetime '`'a`, and `Adventurer<'a>` uses it.

- The lifetime parameter declaration after `impl` and use after the type name is required.

20

Anonymous Lifetime ('_)

- "Anonymous lifetime" is an indicator that references are being used.

```
struct Adventurer<'a> {
    name: &'a str,
    hit_points: u32,
}

impl Adventurer<'_> {
    fn take_damage(&mut self) {
        self.hit_points -= 20;
        println!("{} has {} hit points left!", ... );
    }
}
```

21

Lifetimes and Subtypes

- Lifetimes are just regions of code, and regions can be partially ordered with the *contains* (*outlives*) relationship.
- Subtyping on lifetimes is in terms of that relationship:

if `'big: 'small` ("big *contains* small" or "big *outlives* small"),
then `'big` is a *subtype* of `'small`.

```
struct OutlivesExample<'a, 'b: 'a> {
    a_str: &'a str,
    b_str: &'b str,
}
```

22

Lifetimes and Subtypes (Cont'd)

- If someone wants a reference that lives for '`small`', usually what they actually mean is that they want a reference that lives for at least '`small`'.
- '`static`', the forever lifetime (i.e., the lifetime of the entire program), is a **subtype of every lifetime** because by definition it outlives everything.

```
fn bar<'a>() {
    let s: &'static str = "hi";
    let t: &'a str = s;
}
```

- Since '`static`' outlives the lifetime parameter '`a`', `&'static str` is a subtype of `&'a str`.

23

'static Lifetime

- String literals and globals have the type '`static`' lifetime : they are baked into the data segment of the final binary.

```
let x: &'static str = "Hello, world.";

static FOO: i32 = 5;
let x: &'static i32 = &FOO;
```

- You may also encounter '`static`' as part of a trait bound:

```
fn generic<T>(x: T)
where
    T: 'static {}
```

24

Lifetime Variances

For a type `T<'a>`, '`a` may be:

- **Covariant**: `'b: 'a => T<'b>: T<'a>` (default for immutable data)
- **Contravariant**: `'b: 'a => T<'a>: T<'b>` (only to `fn` parameters)
- **Invariant**: even if `'b: 'a`, nothing can be said about `T<'b>` and `T<'a>`
 1. If the lifetime is present "inside" some sort of `mutable` context
 - whether inside a `&mut` reference, or *interior mutability* like `'cell/RefCell/Mutex`.
 2. If the lifetime is used in multiple spots where the variances conflict.

25

Lifetime and Type Parameter Variances

	<code>'a</code>	<code>T</code>	<code>U</code>
<code>&'a T</code>	covariant	covariant	
<code>&'a mut T</code>	covariant	invariant	
<code>Box<T></code>		covariant	
<code>Vec<T></code>		covariant	
<code>UnsafeCell<T></code>		invariant	
<code>Cell<T></code>		invariant	
<code>fn(T) -> U</code>		contravariant	covariant
<code>*const T</code>		covariant	
<code>*mut T</code>		invariant	

26

What are the variances of each lifetime parameter?

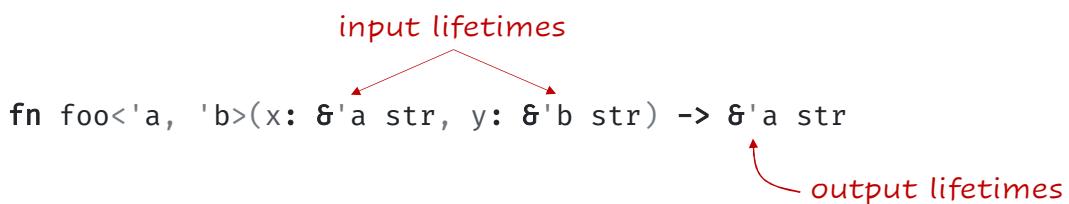
```
struct Multi<'a, 'b, 'c, 'd1, 'd2> {
    a: &'a str,                      // covariant
    b: Cell<&'b str>,              // invariant
    c: fn(&'c str) -> usize,       // contra-variant
    d: &'d1 mut &'d2 str,           // d1: covariant, d2: invariant
}

struct TwoSpots<'a> {           // invariant
    foo: &'a str,
    bar: Cell<&'a str>,
}
```

27

Lifetime Elision

- The patterns programmed into Rust's analysis of references are called the *lifetime elision rules*.
- The elision rules do *not* capture every possible case for lifetimes.
- Lifetimes on function or method parameters are called *input lifetimes*, and lifetimes on return values are called *output lifetimes*.



28

Three rules for elision

1. The compiler assigns a lifetime to each parameter that's a reference.

```
fn foo<'a>(x: &'a i32)  
fn foo<'a, 'b>(x: &'a i32, y: &'b i32)
```

2. If there is exactly one input lifetime, that lifetime is assigned to all output lifetimes.

```
fn foo<'a>(x: &'a i32) -> &'a i32
```

3. If there are multiple input lifetimes, but one of them is `&self` or `&mut self`, the lifetime of `self` is assigned to all output lifetimes.

```
fn foo<'a, 'b>(&'a self, y: &'b i32) -> &'a str
```

29

Lifetimes in Functions and Traits

- Function signatures with lifetimes have a few constraints:
 - any reference must have an annotated lifetime.
 - any reference being returned must have the same lifetime as an input or be `static`.
- Additionally, note that returning references without input is banned if it would result in returning references to invalid data.
- Annotation of lifetimes in trait methods basically are similar to functions. Note that `impl` may have annotation of lifetimes too.

30

Lifetime Coercion

- A *longer lifetime* can be *coerced into* a *shorter one*, so a longer lifetime may be used in place of the shorter lifetime.
- This comes in the form of
 - inferred coercion by the Rust compiler, and also
 - declaring a lifetime difference:

31

```
// Here, Rust infers a lifetime that is as short as possible.  
// The two references are then coerced to that lifetime.  
fn longest<'a>(first: &'a i32, second: &'a i32) -> &'a i32 {  
    if first > second { first } else { second }  
}  
  
// `<'a: 'b, 'b>` reads as lifetime ``a`` is at least as long as ``b``.  
// Here, we take in an `&'a i32` and return a `&'b i32` as a result of coercion.  
fn choose_first<'a: 'b, 'b>(first: &'a i32, _: &'b i32) -> &'b i32 {  
    first  
}  
  
fn fooz() {  
    let first = 2; // Longer lifetime  
  
    {  
        let second = 3; // Shorter lifetime  
  
        println!("The longest is {}", longest(&first, &second));  
        println!("{} is the first", choose_first(&first, &second));  
    };  
}
```

32

Tips to avoid getting too stressed about

- You can stay with owned types, use clones etc. if you want to avoid them for the time being.
- Much of the time, when the compiler wants a lifetime you will just end up writing `<'a>` here and there and then it will work.
 - “*Don't worry, I won't give you anything that doesn't live long enough*”.
- You can explore lifetimes just a bit at a time.
 - Write some code with owned values, then make one a reference.
 - The compiler will start to complain, but also give some suggestions.
 - And if it gets too complicated, you can undo it and try again next time.

33

Lifetimes Summary

- Rust keeps its references under control by assigning lifetimes to them.
- They ensure that types containing references don't outlive their values, which basically prevents us from writing code that produces dangling pointers.
- In many cases, lifetime definitions can be omitted and Rust fills in the gaps for us.
- It's also possible to have types with multiple distinct lifetime parameters.
- Lifetimes are a compile-time only feature and don't exist at runtime.

34