

Concurrency

Mutlithreading

- Improved performance, but added complexity
- Must consider:
 - Race conditions
 - Deadlocks
 - Hard to reproduce and fix bugs, etc.

Creating Threads

- The Rust standard library uses OS threads (1:1 threading).

```
pub fn spawn<F, T>(f: F) -> JoinHandle<T>
where
    F: FnOnce() -> T,
    F: Send + 'static,
    T: Send + 'static,
{
    Builder::new().spawn(f).expect("failed to spawn thread")
}
```

3

Creating Threads (cont'd)

```
use std::thread;
use std::time::Duration;

fn main() {
    thread::spawn(|| {
        for i in 1..10 {
            println!("hi number {} from the spawned thread!", i);
            thread::sleep(Duration::from_millis(1));
        }
    });

    for i in 1..5 {
        println!("hi number {} from the main thread!", i);
        thread::sleep(Duration::from_millis(1));
    }
}
```

4

Waiting for All Threads to Finish Using join Handles

```
fn main() {  
    let handle: JoinHandle<()> = thread::spawn(|| {  
        for i in 1..10 { ... }  
    });  
  
    for i in 1..5 {  
        println!("hi number {i} from the main thread!");  
        thread::sleep(Duration::from_millis(1));  
    }  
  
    handle.join().expect("Couldn't join on the associated thread");  
}  
}
```

thread::spawn returns JoinHandle
waits for its thread to finish
fn join(self) -> Result<T> { ... }

5

What seems to be a problem?

```
let v = vec![1, 2, 3];  
  
let handle = thread::spawn(|| {  
    println!("Here's a vector: {:?}", v);  
});  
  
handle.join().unwrap();
```

Because I'm not going to mutate v inside closure, maybe it's ok to just borrow v.
But ...

- The closure tries to borrow `v` via reference.
- However, Rust can't tell how long the spawned thread will run, so it doesn't know if the reference to `v` will always be valid.

6

```

error[E0373]: closure may outlive the current function, but it borrows `v`,
which is owned by the current function
--> src/main.rs:6:32
   |
6 |     let handle = thread::spawn(|| {
   |                     ^^^ may outlive borrowed value `v`
7 |         println!("Here's a vector: {:?}", v);
   |             - `v` is borrowed here
   |
note: function requires argument type to outlive ``static``
--> src/main.rs:6:18
   |
6 |     let handle = thread::spawn(|| {
   |     ^
7 |     |-----^
8 |     |     println!("Here's a vector: {:?}", v);
   |     |-----^
help: to force the closure to take ownership of `v` (and any other referenced
variables), use the `move` keyword
   |
6 |     let handle = thread::spawn(move || {
   |                     +++

```

7

Using move Closures with Threads

- Adding the `move` keyword before the closure forces the closure to *take ownership of the values* it's using.

```

fn main() {
    let v = vec![1, 2, 3];

    let handle = thread::spawn(move || {
        println!("Here's a vector: {:?}", v);
    });

    handle.join().unwrap();
}

```

8

Why does this code not compile?

```
fn main() {  
    let str = "What's that?";  
  
    let handle = thread::spawn(|| {  
        println!("str = {}", str);  
    });  
  
    handle.join();  
}
```

9

Placing move before closure solves the problem

```
fn main() {  
    let str = "What's that?";  
  
    let handle = thread::spawn(move || {  
        println!("str = {}", str);  
    });  
  
    handle.join();  
}
```

- But, what if we can guarantee that borrowing threads finish before we use the borrowed variables?

10

Scoped Threads (as of Rust 1.63)

- The problem is that when using threads we don't know exactly when the closure we give to `thread::spawn` will run.
- Meaning it is hard to determine when variables borrowed by threads will be dropped. That could lead to *use after free bugs*.

"Scoped threads" ensure threads are joined before the Scope exits.

11

Scoped Thread (cont'd)

- To use scoped threads we first need to create a `Scope` for the scoped threads to operate in.

```
use std::thread::scope;

fn main() {
    let str = "What's that?";

    thread::scope(|s: &Scope| {
        s.spawn(|| {
            println!("str = {}", str);
        });
    });
}
```

12

Scoped threads are guaranteed to be joined before the Scope exits

```
fn guarantee_to_join_before_scope_exits() {  
    let mut i = 0;  
  
    thread::scope(|s| {  
        s.spawn(|| {  
            thread::sleep(Duration::from_millis(1000));  
            i += 1;  
        });  
    }); // thread is joined here  
  
    println!("i = {}", i); // it's safe to use variable i  
}
```

13

Shared Mutable State

"A widely accepted method of communication is by inspection and updating of a common store ... However, this can create severe problems in the construction of correct programs and it may lead to expense and ... unreliability (e.g., glitches) ..."

– Tony Hoare, CSP 1978



"Don't communicate by sharing memory;
instead,
share memory by communicating"



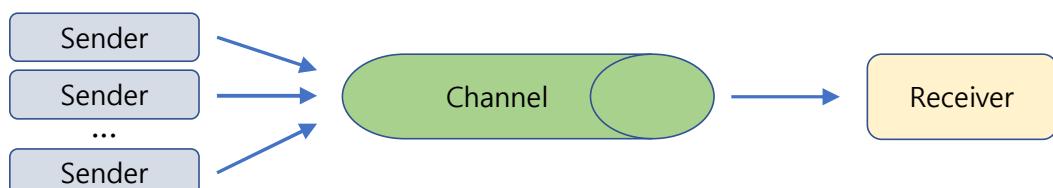
14

mpsc Channels

- Threads can communicate with each other via *channels*.



- Multiple producers and a single consumer can be associated with a channel (mpsc model).



15

Creating Channels

- Create a new channel using the `mpsc::channel` function.
 - `mpsc` stands for multiple producer, single consumer.
- A channel is said to be *closed* if either the transmitter or receiver half is dropped.

```
use std::sync::mpsc;

fn main() {
    let (tx, rx) = mpsc::channel();
}

transmitter   receiver
```

16

Single Producer Single Receiver

```
use std::sync::mpsc::{ self, Sender, Receiver };
use std::thread;

fn main() {
    let (tx, rx): (Sender<String>, Receiver<String>) = mpsc::channel();

    thread::spawn(move || {
        let val = String::from("hi");
        tx.send(val).unwrap();
    });

    let received = rx.recv().unwrap();
    println!("Got: {received}");
}
```

17

```
error[E0382]: borrow of moved value: `val`
--> src/main.rs:10:31
|
8 |         let val = String::from("hi");
|             --- move occurs because `val` has type `String`, which does not implement the
`Copy` trait
9 |         tx.send(val).unwrap();
|             --- value moved here
10|         println!("val is {}", val);
|                         ^^^ value borrowed here after move
|
```

```
thread::spawn(move || {
    let val = String::from("hi");
    tx.send(val).unwrap();
    println!("val is {val}");
});

let received = rx.recv().unwrap();
println!("Got: {received}");
}
```

val was already moved ☺

18

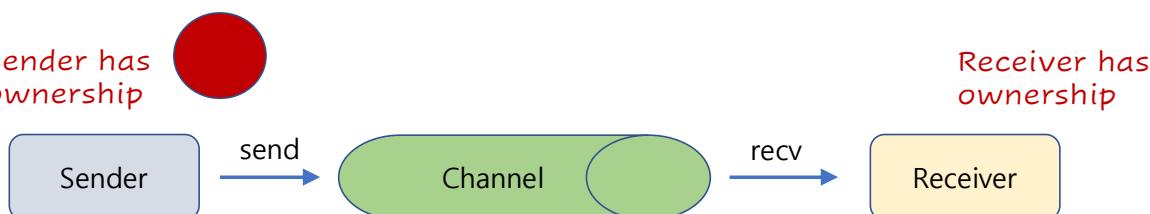
Channels and Ownership Transference

- The ownership rules play a vital role in message sending because they help you write safe, concurrent code.
- The `send` function takes ownership of its parameter, and when the value is moved, the receiver takes ownership of it.
- This stops us from accidentally using the value again after sending it; the ownership system checks that everything is okay.

19

Channels and Ownership Transference

- The ownership rules play a vital role in message sending because they help you write safe, concurrent code.



- This stops us from accidentally using the value again after sending it; the ownership system checks that everything is okay.

20

Creating Multiple Producers by Cloning the Transmitter

```
let (tx, rx) = mpsc::channel();
let tx1 = tx.clone();
thread::spawn(move || {
    let vals = vec![String::from("hi"), ... , String::from("thread")];
    for val in vals {
        tx1.send(val).unwrap(); thread::sleep(Duration::from_secs(1));
    }
});
thread::spawn(move || {
    let vals = vec![String::from("more"), ... ,String::from("you")];
    for val in vals {
        tx.send(val).unwrap(); thread::sleep(Duration::from_secs(1));
    }
});
for received in rx { println!("Got: {received}"); }
```

21

Channel vs. sync_channel

- The `recv` method for both channel types block until a message is received.
- In Channel, calls to `send` do not block the sending thread.

```
let (sender, receiver) = mpsc::channel();
sender.send(2); // Does not block, even if no receiver handles value
```

- A sender created with `sync_channel` may block. When we call `sync_channel` we specify a queue size. If the queue is full a call to `send` will block until space on the queue is available.

```
use std::sync::mpsc::{ self, Sender, Receiver };
```

```
let (sender, receiver): (Sender<i32>, Receiver<i32>) = mpsc::channel();
let (sender, receiver) = mpsc::sync_channel::<String>(1); // queue size
```

22

Shared State Concurrency

- While channels are similar to single ownership, shared memory concurrency is like multiple ownership.
- Smart pointers made multiple ownership possible.
- But, multiple ownership can add complexity because these different owners need managing.
- Management of mutexes can be incredibly tricky to get right.
- Rust's **type system and ownership rules** greatly assist in getting this management correct.

23

Data Races

A data race is defined to occur when two distinct threads access the same memory location, where

- at least one of them is a write, and
- there is no synchronization mechanism that enforces an ordering on the accesses.

24

Using Mutexes to Allow Access to Shared Data

- A mutual exclusion primitive useful for protecting shared data
- Unlike C++, the value we want to protect is literally inside the mutex.
 - The **mutex** is described as **guarding the data** it holds via the locking system.
- The data can only be accessed through **MutexGuard** returned from **lock** and **try_lock**.

```
pub struct Mutex<T: ?Sized> {    pub struct MutexGuard<'a, T: ?Sized + 'a> {  
    inner: sys::Mutex,  
    poison: poison::Flag,  
    data: UnsafeCell<T>,        lock: &'a Mutex<T>,  
    }                                poison: poison::Guard,  
}  
}
```

25

The API of Mutex<T>

```
// Create a mutex that holds an i32.  
let m: Mutex<i32> = Mutex::new(5);  
  
{  
    // To access data inside mutex, use `lock` to acquire the lock.  
    // This call will block the current thread until it has the lock.  
    // The call to lock returns a smart pointer called `MutexGuard`, wrapped  
    // in a `LockResult` that we handled with the call to `unwrap`.  
    let mut num: MutexGuard<i32> = m.lock().unwrap();  
  
    // `MutexGuard` smart pointer implements `Deref` to point at our inner data;  
    *num = 6;  
} // The smart pointer also has a `Drop` implementation that releases the lock  
// automatically when a `MutexGuard` goes out of scope  
  
println!("m = {m:?}");  
}
```

26

C++ vs. Rust

```

int cplus_mutex() {
    int counter = 0;
    std::mutex mtx;
    std::vector<std::thread> handles;
    for (int j = 0; j < 10; ++j) {
        handles.push_back(
            std::thread([&mtx, &counter]() {
                std::scoped_lock lk(mtx);
                counter += 1; ←
            }));
    }                                protect codes
    for (auto &handle : handles) {
        handle.join();
    }
    std::cout << counter;
}

```

protect data

```

fn rust_mutex() {
    let counter = Arc::new(Mutex::new(0));
    let mut handles = vec![];
    for _ in 0..10 {
        let counter = Arc::clone(&counter);
        let handle = thread::spawn(move || {
            let mut num = counter.lock().unwrap();
            *num += 1;
        });
        handles.push(handle);
    }
    for handle in handles {
        handle.join().unwrap();
    }
    println!("{}", *counter.lock().unwrap());
}

```

What's wrong?

```

// Wrong ver. 1
let counter = Mutex::new(0);
let mut handles = vec![];
for _ in 0..10 {
    let handle = thread::spawn(move || {
        let mut num = counter.lock().unwrap();
        *num += 1;
    });
    handles.push(handle);
}

```

value moved into closure here,
in previous iteration of loop

Rc<T> cannot be used across
thread boundaries since
reference counting operation
is not atomic.

// Wrong ver. 2

```

let counter = Rc::new(Mutex::new(0));
let mut handles = vec![];

```

```

for _ in 0..10 {
    let counter = Rc::clone(&counter);
    let handle = thread::spawn(move || {
        let mut num = counter.lock().unwrap();
        *num += 1;
    });
    handles.push(handle);
}

```

Arc<Mutex<T>> for Multithreading



- `std::sync::Arc<T>` is an *atomically reference counted smart pointer* type, i.e., thread-safe equivalent of `Rc<T>`.
- `std::sync::Mutex<T>` provides *interior mutability*, as the `Cell`/`RefCell` do.
- Use `Rc<RefCell<T>>` in sequential setting
- Use `Arc<Mutex<T>>` in concurrent setting

```
use std::sync::{Arc, Mutex};

fn share_mutex_between_threads() {
    let counter = Arc::new(Mutex::new(0));
    let mut handles = vec![];

    for _ in 0..10 {
        let counter = Arc::clone(&counter);
        let handle = thread::spawn(move || {
            let mut num = counter.lock().unwrap();
            *num += 1;
        });
        handles.push(handle);
    }

    for handle in handles {
        handle.join().unwrap();
    }

    println!("{}", *counter.lock().unwrap());
}
```

29

Extensible Concurrency with the Sync and Send Traits

- Rust language has very few concurrency features.
- Almost every concurrency feature we've talked about so far has been *part of the standard library*, not the language.
- You can write your own concurrency features or use those written by others, if necessary.
- However, two concurrency concepts are embedded in the language: the `std::marker::{Sync, Send}` traits.
- Implementing `Send` and `Sync` manually is unsafe.

Trait std::marker::Send

- Types that can be transferred across thread boundaries.
- Send types
 - most primitive types,
 - [Arc](#)
 - [Cell](#), [Mutex](#)
- Non-Send types
 - [Rc](#), [MutexGuard](#)

```
pub unsafe auto trait Send { }
```

31

Allowing Transference of Ownership Between Threads with Send

- The [Send](#) marker trait indicates that **ownership** of values of the type implementing [Send](#) **can be transferred** between threads.
- Almost every Rust type is [Send](#), but there are some exceptions, including [Rc<T>](#).

```
impl<T: ?Sized> !Send for Rc<T> {}
impl<T: ?Sized> !Send for MutexGuard<'_, T> {}
```

- Almost all primitive types are [Send](#) aside from raw pointers.

```
impl<T: ?Sized> !Send for *const T {}
impl<T: ?Sized> !Send for *mut T {}
```

- Any type composed entirely of [Send](#) types is automatically marked as [Send](#) as well.

32

Trait std::marker::Sync

- Types for which it is safe to share references between threads.

A type `T` is `Sync` if and only if `&T` is `Send`.

- Sync types

- Primitive types like `u8`, `f64` etc. and simple aggregate types like tuples, enums, and structs containing them
- `&T`, `&mut T`, `Box<T>`, `Vec<T>` and most other collection types (if `T` is `Sync`)
- Generic parameters need to be `Sync` for their container to be `Sync`
- `Arc`, `Mutex` and `RwLock`, and atomic data types

- Non-Sync types

- `Rc`, `Cell`, `RefCell`

33

Allowing Access from Multiple Threads with Sync

- The `Sync` indicates that it is safe for the type implementing `Sync` to be referenced from multiple threads.
- In other words, any type `T` is `Sync` if `&T` (an immutable reference to `T`) is `Send`, meaning the reference can be sent safely to another thread.
- Similar to `Send`, primitive types are `Sync`, and types composed entirely of types that are `Sync` are also `Sync`.
- `Rc<T>`, `RefCell<T>`, and the family of related `Cell<T>` types are not `Sync`.
- `Mutex<T>` is `Sync`.

34

Summary

- The Rust standard library provides channels for message passing and smart pointer types, such as `Mutex<T>` and `Arc<T>`, that are safe to use in concurrent contexts.
- The type system and the borrow checker ensure that the code using these solutions won't end up with data races or invalid references.
- Once you get your code to compile, you can rest assured that it will happily run on multiple threads without the kinds of hard-to-track-down bugs common in other languages.
- Concurrent programming is no longer a concept to be afraid of: go forth and **make your programs concurrent, fearlessly!**

35

36