



## Futures and async/await

1

## Asynchronous Programming

- Callbacks
- Multi-threading
- Future/Promise/Rx
- Coroutines

- Responsiveness
- Performance



2

# Essentials

An async application should pull in at least two crates from Rust's ecosystem:

1. **futures**, an official Rust crate that lives in the rust-lang repository
2. A runtime of your choosing, such as **Tokio**, **async\_std**, **smol**, etc.

```
[dependencies]
futures = { version = "0.3.*" }
tokio = { version = "1.24.*", features = [ "full" ] }
async-std = {
    version = "1.12.0",
    features = [ "unstable" ], [ "attributes" ]
}
```

3

# Futures

- Futures represent values that *should exist at some point in the future* as the result of an async operation.
- Concurrency primitives similar to **Promises** in Javascript.
- However, **Rust's Future is lazy** in that it does nothing until explicitly requested via such as **await**.
- Stepping stone to **async/await** which allows users to write *asynchronous code that looks like synchronous code*.
- Futures arrived in stable Rust with version 1.36.0.
- Rust is a low-level language and doesn't include a runtime for scheduling async tasks.

4

## async fn() syntax

- Use `async` in front of `fn`, `closure`, or a `block` to turn the code into a `Future`.
- `async` is a keyword to annotate an *async function*

```
async fn foo() -> u32 {  
    42  
}
```

which is just a syntactic sugar for:

```
fn foo_desugared -> impl Future<Output = u32> {  
    async { 42 }  
}
```

```
let f: impl Future<Output = u32> = foo();  
let f: impl Future<Output = u32> = foo_desugared();  
let f: impl Future<Output = u32> = async { 42 };
```

5

## async/.await

- `async/.await` is Rust's built-in tool for *writing asynchronous functions that look like synchronous code*.
- Similar to `async/await` in Javascript and C#, etc.
  - Unlike Javascript, you need to pick a runtime to actually run your asynchronous code.
- `async` is zero-cost. You can use `async` without heap allocations and dynamic dispatch.
- `async/.await` is a simpler syntax for futures.

6

## async/.await (cont'd)

- Async functions will not be run immediately, but will only be evaluated when the returned future is submitted to the executor via `.await` etc.

```
async fn hello_rusty() {  
    println!("Hello, Rusty!");  
}  
  
async fn execute_async_function() {  
    let future = hello_rusty(); // Nothing is printed  
    future.await; // `future` is run and "hello, world!" is printed  
}  
    ↗ .await is only allowed inside async functions and blocks
```

- `.await` yields control back to the calling thread to make other tasks progress.
- Rust does not let you declare `async` functions in traits, yet.

7

## Runs a Future

- `block_on` run a future to completion on the current thread.
  - This function will block the caller until the given future has completed.

```
async fn hello_rusty() {  
    println!("Hello, Rusty!");  
}
```

```
use futures::executor;  
executor::block_on(hello_rusty());  
// "hello, world!" is printed
```

```
use async_std::task;  
task::block_on(hello_rusty());
```

```
use tokio::runtime;  
let rt = runtime::Runtime::new().unwrap();  
rt.block_on(hello_rusty());
```

8

# From Synchronous to Asynchronous

```
fn postItem(item: Item) {  
    let token = requestToken();  
    let post = createPost(token, item);  
    showPost(post)  
}  
  
fn requestToken() -> Token {  
    // makes request for a token & waits  
    token // returns result when received  
}  
  
fn createPost(token: Token, item: Item) -> Post {  
    // sends item to the server & waits  
    return post // returns resulting post  
}  
  
fn showPost(post: Post) {  
    // does some local processing of result  
}
```

9

## Async via Callbacks


```
fn post_item(item: Item) {  
    request_token(move |token| {  
        create_post(token, item, move |post| {  
            show_post(post);  
        });  
    });  
}  
  
fn request_token<F>(callback: F) where F: Fn(Token) + Send + 'static {  
    background(Token, callback);  
}  
  
fn create_post<F>(token: Token, item: Item, callback: F)  
where F: Fn(Post) + Send + 'static {  
    background(Post, callback);  
}  
  
fn show_post(post: Post) {  
    println!("post displayed");  
}
```

👎 Hard to read and  
harder to reason about

👎 Handling exceptions  
makes it a real mess

👎 Callback Hell

10



```

private fun loadData() {
    networkRequest { data ->
        anotherRequest(data) { data1 ->
            anotherRequest(data1) { data2 ->
                anotherRequest(data2) { data3 ->
                    anotherRequest(data3) { data4 ->
                        anotherRequest(data4) { data5 ->
                            anotherRequest(data5) { data6 ->
                                anotherRequest(data6) { data7 ->
                                    anotherRequest(data7) { data8 ->
                                        anotherRequest(data8) { data9 ->
                                            anotherRequest(data9) {
                                                // How many more do you want?
                                                println(it)
                                            }
                                        }
                                    }
                                }
                            }
                        }
                    }
                }
            }
        }
    }
}

```

**Callback Hell**

11

## Async via Multi-threads

```

fn post_item(item: Item) {
    let (tx, rx) = std::sync::mpsc::channel();

    std::thread::spawn(move || {
        tx.send(request_token()).unwrap();
    });

    let (tx2, rx2) = std::sync::mpsc::channel();
    std::thread::spawn(move || {
        let token = rx.recv().unwrap();
        tx2.send(create_post(token, item)).unwrap();
    });

    std::thread::spawn(move || {
        let post = rx2.recv().unwrap();
        show_post(post);
    });
}

```

 Spawning threads can be expensive

 Codes can get quite complicated

 Error-prone

12

# Async vis Future Combinators

```
fn post_item(item: Item) -> impl Future<Output = ()> {
    request_token()
        .then(move |token| create_post(token, item))
        .then(|post| {
            show_post(post);
            futures::future::ready(())
        })
}
}

impl Future for RequestToken {
    type Output = Token;

    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
        // Some codes here
        Poll::Ready(Token)
    }
}
```

-  No nesting indentation
-  Composable
-  Library-specific operators
-  Manually implemented futures

13

# Async via `async/.await`

```
async fn post_item(item: Item) {
    let token = request_token().await;           // suspension point
    let post = create_post(token, item).await; // suspension point
    show_post(post);
}

async fn request_token() -> Token {
    delay(2000); // simulate network delay
    Token
}

async fn create_post(token: Token, item: Item) -> Post {
    delay(2000); // simulate network delay
    Post
}

fn show_post(post: Post) {
    println!("post displayed");
}
```

*The async world is nicely sequential!*

14

# Concurrency using async/.await

```
async fn loadImage(name: String): Image = { ... }
fn combineImages(img1: &Image, &img2: Image): Image = { ... }
```

```
// Non-blocking, but sequential code. How to make this code run concurrently?
async fn load_and_combine(name1: String, name2: String): Image {
    let task1 = loadImage(name1);
    let task2 = loadImage(name2);

    combineImages(task1.await, task2.await)
}
```



suspends until task is complete

15

# Concurrency using async/.await

```
async fn loadImage(name: String): Image = { ... }
fn combineImages(img1: &Image, &img2: Image): Image = { ... }
```

```
// Non-blocking, but sequential code. How to make this code run concurrently?
async fn load_and_combine(name1: String, name2: String): Image {
    let task1 = loadImage(name1);
    let task2 = loadImage(name2);

    combineImages(task1.await, task2.await)
}
```



suspends until task is complete

16

# Three Basic Operations to Figure Out

- How to start the runtime
- How to spawn a Future
- How to spawn blocking or CPU-intensive tasks

We will choose **tokio** as our runtime.

17

## Async runtime - tokio

Tokio provides:

- Top-level ‘loop’
- Low-level resources like network sockets and timers (e.g.,  
`tokio::net::TcpListener`, `tokio::fs`, `tokio::timer`)
- **Executor** schedules and drives futures to completion.
- **Reactor** interacts with OS (via crate Mio) to say ‘Wake me up in any of these file descriptors change state’ (epoll, kqueue , IO Completion Ports).
- Also handles waiting on non-OS events, e.g., receiving on channels.

18

# Starting the runtime

- Explicitly instantiate the runtime, or use attribute macro `#[tokio::main]`.

```
use tokio::runtime::Runtime;                                #[tokio::main]
fn main() {                                                 async fn main() {
    let mut rt = Runtime::new().unwrap();                   app().await;
    let future = app();                                    }
    rt.block_on(future);                                 fn main() {
}                                                       tokio::runtime::Builder
async fn app() {                                         ::new_multi_thread()
    println!("Hello, Tokio!");                           .enable_all()
}                                                       .build()
                                                       .unwrap()
                                                       .block_on(app())
}
```



19

# Spawning a Future on the runtime

```
use std::time::Duration;
use tokio::time::sleep;

#[tokio::test]
async fn spawning_task() {
    use tokio::task;

    async fn some_task() {
        sleep(Duration::from_millis(1000)).await;
        println!("Hello from some_task");
    }

    let future = task::spawn(some_task()); // starts running now
    let result = future.await; // no need to await, if not necessary
}
```

20

# Spawning blocking or CPU-intensive tasks

- You should avoid blocking or running CPU-intensive code in Futures themselves.
- Most of the code you write in async Rust will actually be executed in a Future.
- You should offload this work to a different thread, so that you avoid blocking the executor thread.
- In tokio, you can do this via `task::spawn_blocking`.

```
fn fibonacci(n: u32) -> u32 {
    match n {
        0 | 1 => n,
        v => fibonacci(v - 1) + fibonacci(v - 2),
    }
}
```

21

# Spawning blocking or CPU-intensive tasks (cont's)

```
async fn foo() {
    for i in 0..10 {
        println!("Hello from foo: {i}");
        sleep(Duration::from_millis(500)).await;
    }
}
```

```
let fib = async { fibonacci(45) };
let (result, _) = tokio::join!(fib, foo());
println!("Result: {result}");
```



```
let fib = task::spawn_blocking(|| fibonacci(45));
let (result, _) = tokio::join!(fib, foo());
println!("Result: {}", result.unwrap());
```



22

# Create futures Runtime (futures::future::\*)

- Provides abstractions for asynchronous programming: Futures, Streams, Sinks, and Executors.
- Also provides a lot combinators (`map`, `then`, etc.) for future composition.
- `block_on` run a future to completion on the current thread.
  - This function will block the caller until the given future has completed.

```
use std::future::Future;
use futures::{prelude::*, executor::block_on}

// Creates a future that is immediately ready with a value.
let future = future::ready(42);
let result = block_on(future);

// Create a future that is immediately ready with a success value.
let future = future::ok::<i32, &str>(42);

// Create a future that is immediately ready with an error value.
let future = future::err::<i32, &str>("error");
```

23

# Futures in Depth

24


# Trait std::future::Future

- Future is polled by an executor as many time as necessary until its value is ready.

```
trait Future {  
    type Output;  
  
    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output>;  
}
```

Essentially a Callback to be called when the future is ready to progress.

```
enum Poll<T> {  
    Ready(T),  
    Pending,  
}
```



25

# Executors and Runtime

- Runtimes normally consist of **Executor** and **Reactor**.
- Executor creates a *waker* and passes it to the Future when it is time to poll.
- Executor keeps invoking polling method of the Future whenever the passed *waker* notifies it, which will drive the future to completion..
- Reactor notifies the Executor of event completion via *waker*.
- There are no built-in runtime including executors in the standard library.
- Instead, runtimes are provided by community maintained crates (e.g., **tokio**, **async-std**, **smol**, etc.).
  - This allows different applications to choose different executors to make trade-offs that best meet the goals of a given applications.
  - Both single- and multi-threaded runtimes are available, which have different strengths and weaknesses.

26

# When does the task run?

- A **task** is a top-level future that has been submitted to an executor, and it is a unit of asynchronous computation.
- Some executor at a top level polls futures until progress can be made.
  - A very basic executor might poll futures in a loop until they complete.

```
Future {  
    processing  
    Future.await  
    processing  
    Future.await  
    processing  
}
```

Future



27

# Simplified Top-level Executor


```
pub fn run_block_on<F: Future>(future: F) -> F::Output {  
    pin_mut!(future); // Pins a value on the stack  
  
    loop {  
        // Create a waker  
        let t = std::thread::current();  
        let waker = Arc::new(ThreadWaker(t)).into();  
        let mut cx = Context::from_waker(&waker);  
  
        match future.as_mut().poll(&mut cx) {  
            Poll::Ready(output) => return output,  
            Poll::Pending => {  
                std::thread::park();  
            }  
        }  
    }  
}
```

28

# The .await expands roughly to this:

```
let mut pin = unsafe { Pin::new_unchecked(&mut future) };
loop {
    match Future::poll(&mut pin, &mut ctx) {
        Poll::Ready(item) => Poll::Ready(item),
        Poll::Pending      => Poll::Pending, // yield
    }
}
```

29



30

# Implicitly Generated Future as State Machine

```
async fn foo() {
    println!("{}:?", starts);
    let result = io_read().await;
    println!("{}:?", result);
}

enum Status<F>
where
    F: Future<Output = [u8; 10]>,
{
    Step0,
    Step1(F),
    Step2,
}

```

```
struct Foo {
    status: Status<IoRead>,
}

impl Foo {
    fn new() -> Self {
        Self {
            status: Status::Step0,
        }
    }
}

async fn io_read() -> [u8; 10] {
    let mut buffer = [0u8; 10];
    let mut b = "This string will be read"
        .as_bytes();
    // read up to 10 bytes
    b.read(&mut buffer);
    buffer
}
```

31

# Implicitly Generated Future as State Machine

```
impl Future for Foo {
    type Output = ();

    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
        let mut status = unsafe {
            &mut self.get_unchecked_mut().status
        };

        match status {
            Status::Step0 => {
                println!("{}:?", starts);
                *status = Status::Step1(IoRead::new());
                cx.waker().wake_by_ref();
                Poll::Pending
            }
        }
    }
}
```

32

```
    Status::Step1(ref mut ioread) => {
        match Pin::new(ioread).poll(cx) {
            Poll::Ready(result) => {
                *status = Status::Step2(result);
                cx.waker().wake_by_ref();
                Poll::Pending
            }
            Poll::Pending => {
                cx.waker().wake_by_ref();
                Poll::Pending
            }
        }
    }
    Status::Step2(result) => {
        println!("{:?}", result);
        Poll::Ready(())
    }
}
```

33

# Self-referential Types

```

struct Test {
    a: String,
    b: *const String,
}

impl Test {
    fn new(txt: &str) -> Self {
        Test {
            a: String::from(txt),
            b: std::ptr::null(),
        }
    }

    fn init(&mut self) {
        self.b = &self.a as *const String;
    }

    fn a(&self) -> &str { &self.a }
    fn b(&self) -> &String { unsafe { &*(self.b) } }
}


let mut test1 = Test::new("test1");
test1.init();
let mut test2 = Test::new("test2");
test2.init();

println!("a: {}, b: {}", test1.a(), test1.b());
println!("a: {}, b: {}", test2.a(), test2.b());
// Use the swap() function to swap the two
std::mem::swap(&mut test1, &mut test2);
println!("a: {}, b: {}", test1.a(), test1.b());
println!("a: {}, b: {}", test2.a(), test2.b());

```

34

# Self-referential Types



35

# Pinning and boxing

- Rust does not know whether a type can be safely moved.
- Types that shouldn't be moved must be wrapped inside `Pin<T>`.
- Most types are `Unpinned` types. They implement the trait `Unpin` and can be freely moved within memory.
- If a type is wrapped inside `Pin<T>` and the wrapped type is `!Unpin`, it is not possible to get a mutable reference out of it.
- Futures created by the `async` keyword are `!Unpin` and thus must be pinned.

36

# Implicitly Generated Future as State Machine

```
async fn foo() {
    println!("{}:?", starts);
    let mut buffer = [u8; 10];
    io_read(&mut buffer).await;
    println!("{}:?", buffer);
}

enum Status<F>
where
    F: Future<Output = ()>,
{
    Step0,
    Step1(F),
    Step2,
}
```

```
struct Foo {
    buffer: [u8; 10],
    status: Status<IoRead>,
}
impl Foo {
    fn new() -> Self {
        Self {
            buffer: [0u8; 10],
            status: Status::Step0,
        }
    }
}
async fn io_read(mut buf: &mut [u8; 10]) {
    let mut b = "This string will be read"
        .as_bytes();
    // read up to 10 bytes
    b.read(&mut buffer[..]);
}
```

37

# Implicitly Generated Future as State Machine

```
impl Future for Foo {
    type Output = ();

    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
        let mut future = unsafe { self.get_unchecked_mut() };

        match future.status {
            Status::Step0 => {
                println!("{}:?", starts);
                status = Status::Step1(IoRead::new(&mut future.buffer as *mut_));
                cx.waker().wake_by_ref();
                Poll::Pending
            }
        }
    }
}
```

38

```
    Status::Step1(ref mut ioread) => {
        match Pin::new_unchecked(ioread).poll(cx) {
            Poll::Ready(_) => {
                *status = Status::Step2;
                cx.waker().wake_by_ref();
                Poll::Pending
            }
            Poll::Pending => {
                cx.waker().wake_by_ref();
                Poll::Pending
            }
        }
    }
    Status::Step2 => {
        println!("{}:{}", future.buffer);
        Poll::Ready(())
    }
}
```

39

40